



FACULDADE DE  
CIÊNCIAS E TECNOLOGIA  
DEPARTAMENTO DE INFORMÁTICA

# The Relative Power of Synchronization Operations

lectures 11 & 12 (2025-04-14)

**Master in Computer Science and Engineering**

— Concurrency and Parallelism / 2024-25 —


João Lourenço <[joao.lourenco@fct.unl.pt](mailto:joao.lourenco@fct.unl.pt)>

Based in the companion slides from “The Art of Multiprocessor Programming”

# Outline

- Concurrent Objects
  - Correctness
  - Sequential Objects
  - Quiescent and Sequential Consistency
  - Linearizability
  - Progress Conditions
- Bibliography:
  - **Chapters 5** of book

Herlihy M., Shavit N., Luchangco V., Spear M.;  
**The Art of Multiprocessor Programming**;  
Morgan Kaufmann (2020); ISBN: 978-0-12-415950-1




# Last Lecture

---



# Shared-Memory Computability?



- Mathematical model of **concurrent** computation
- What is (and is not) **concurrently** computable
- Efficiency (mostly) irrelevant

# Wait-Free Implementation


- Every method call completes in finite number of steps
- Implies no mutual exclusion



# From Weakest Register

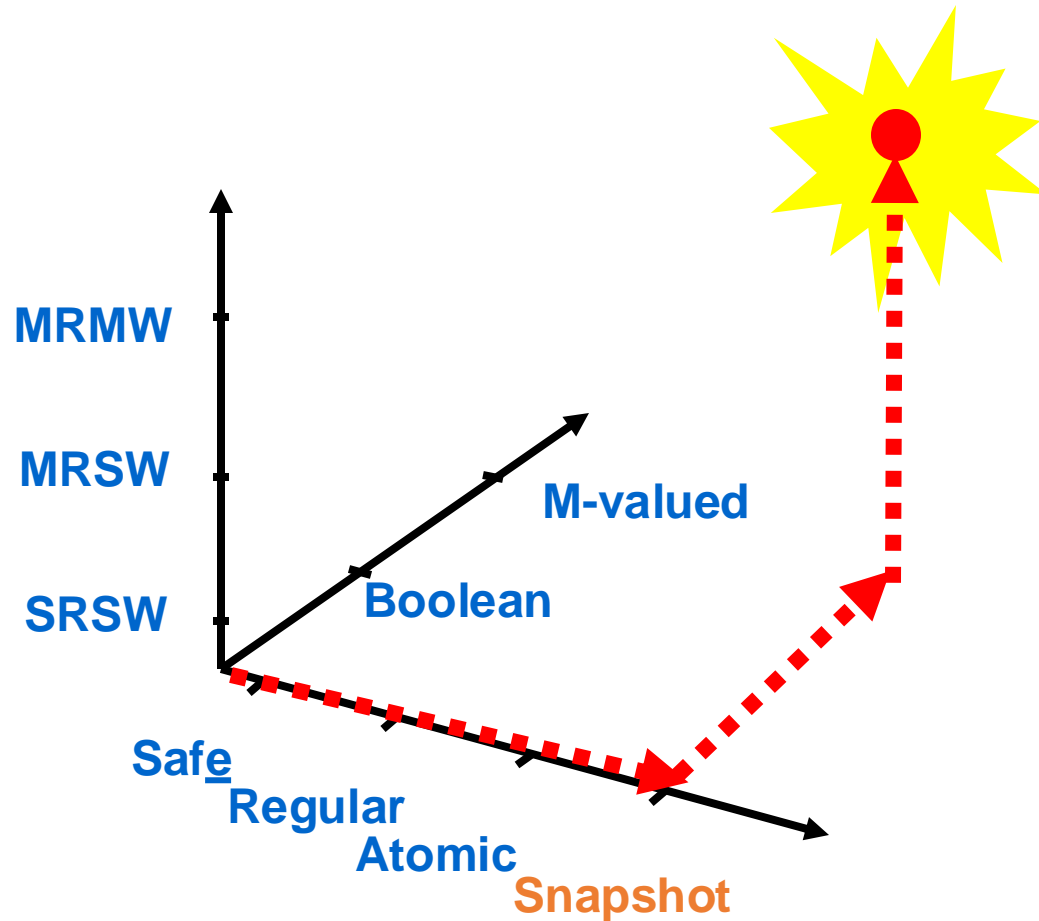
Single writer

Single reader



Safe Boolean register

# All the way to a Wait-free Implementation of Atomic Snapshots



# Rationale for wait-freedom

---

- We wanted atomic registers to implement mutual exclusion



# Rationale for wait-freedom

---

- We wanted atomic registers to implement mutual exclusion
- So we couldn't use mutual exclusion to implement atomic registers

# Rationale for wait-freedom

---

- We wanted atomic registers to implement mutual exclusion
- So we couldn't use mutual exclusion to implement atomic registers
- But wait, there's more!

# This Lecture

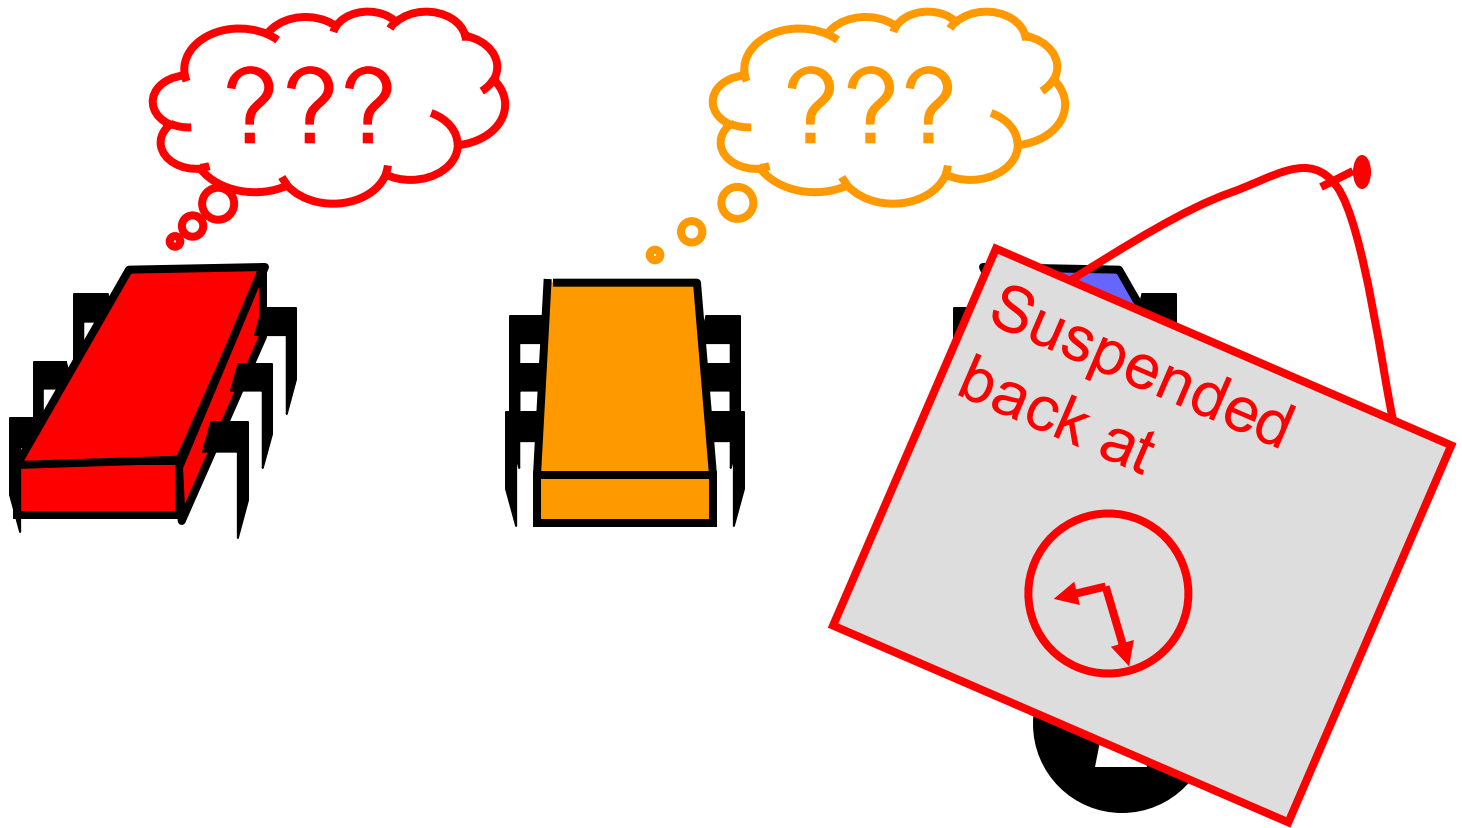
---




---

# Why is Mutual Exclusion so wrong?


# Asynchronous Interrupts



# Heterogeneous Processors



# Fault-tolerance



# Basic Questions

---

- Wait-Free synchronization might be a good idea in principle




# Basic Questions


---

- Wait-Free synchronization might be a good idea in principle
- But how do you do it ...
  - Systematically?
  - Correctly?
  - Efficiently?

# FIFO Queue: Enqueue Method




# FIFO Queue: Dequeue Method




# Two-Thread Wait-Free Queue


```
public class WaitFreeQueue {  
    int head = 0, tail = 0;  
    Item[QSIZE] items;  
    public void enq(Item x) {  
        while (tail-head == QSIZE) {};  
        items[tail % QSIZE] = x; tail++;  
    }  
    public Item deq() {  
        while (tail-head == 0) {}  
        Item item = items[head % QSIZE];  
        head++; return item;  
    }  
}
```



# What About Multiple Dequeueers?



# Grand Challenge



Only new  
aspect

- Implement a FIFO queue
  - Wait-free
  - Linearizable
  - From atomic read-write registers
  - Multiple dequeuers

# Puzzle


---

While you are ruminating on the grand challenge ...

We will give you another puzzle ...


Consensus!

# Consensus: Each Thread has a Private Input







# They Communicate



# They Communicate



# They Agree on One Thread's Input



# Formally: Consensus

---


- Consistent:
  - all threads decide the same value

# Formally: Consensus

---

- Consistent:
  - all threads decide the same value
- Valid:
  - the common decision value is some thread's input

# No Wait-Free Implementation of Consensus using Registers



# Formally

---

- Theorem
  - There is no wait-free implementation of  $n$ -thread consensus from read-write registers

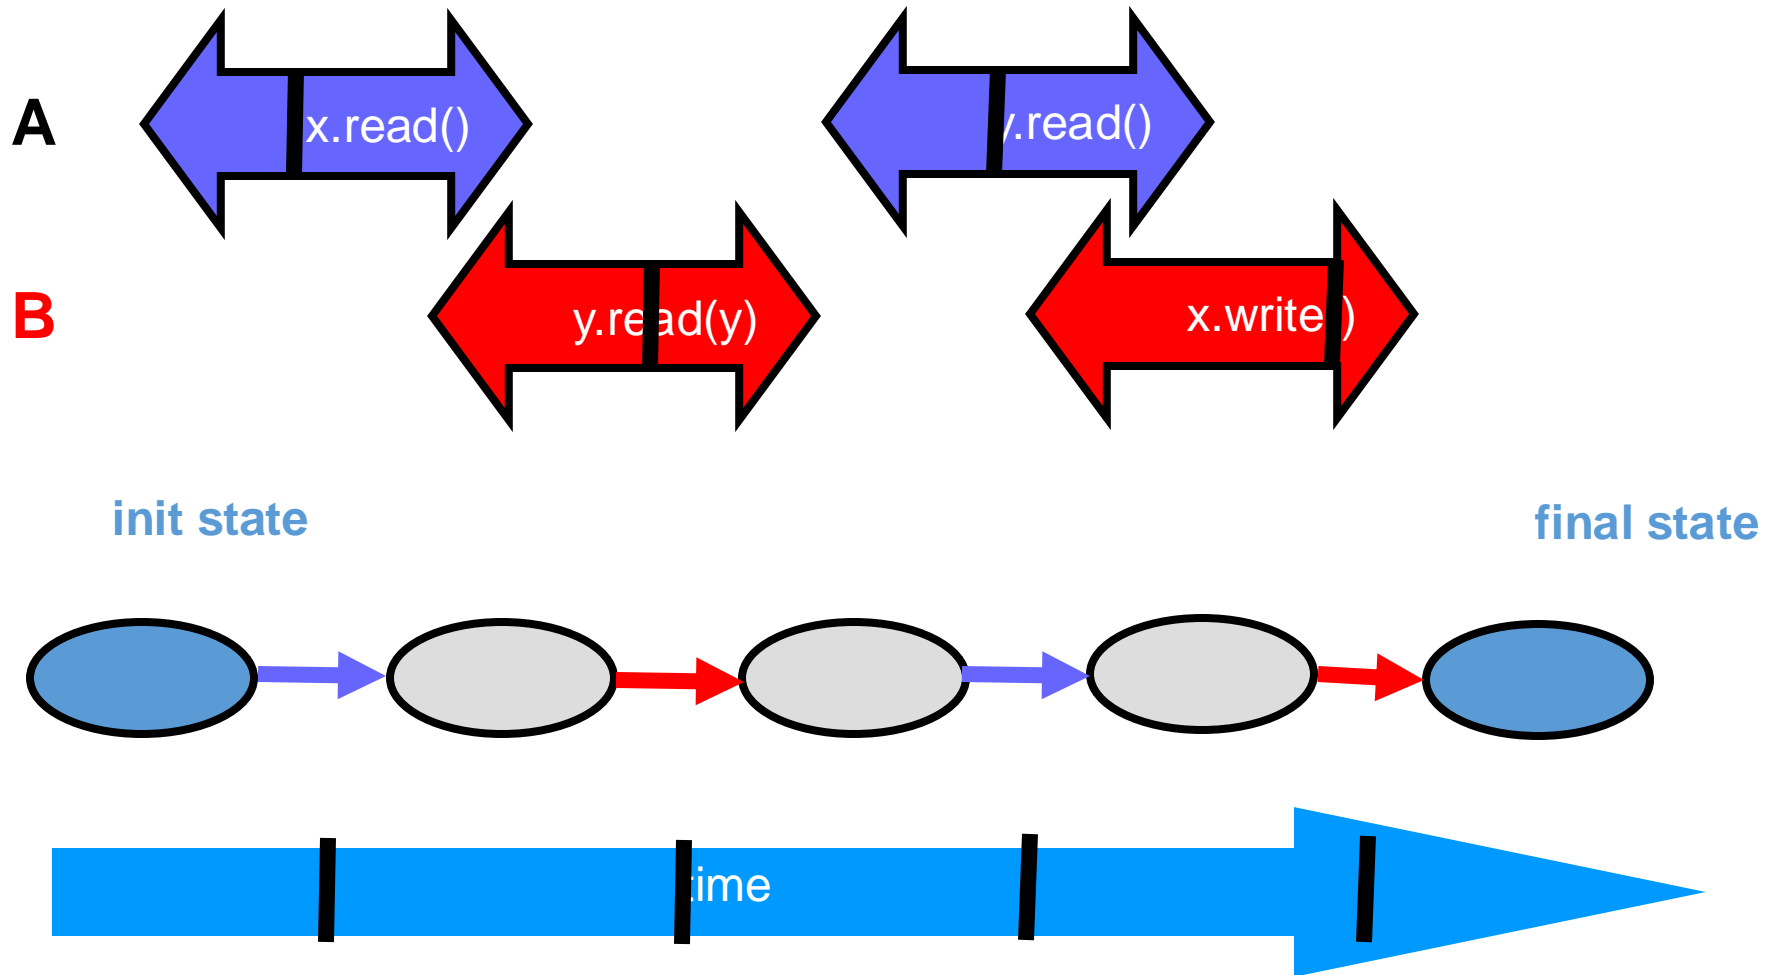
# Formally

---


- Theorem
  - There is no wait-free implementation of  $n$ -thread consensus from read-write registers
- Implication
  - Asynchronous computability  $\neq$  sequential (Turing) computability
- Proof
  - Read the book chapter 5



# Protocol Histories as State Transitions




# Wait-Free Computation




- Either A or B “moves”
- Moving means
  - Register read
  - Register write


# The Two-Move Tree




# Decision Values




# Bivalent: Both Possible



# Univalent: Single Value Possible



# x-valent: x Only Possible Decision



# Summary

---

- Wait-free computation is a tree



# Summary

---

- Wait-free computation is a tree
- Bivalent system states
  - Outcome not fixed

# Summary

---

- Wait-free computation is a tree
- Bivalent system states
  - Outcome not fixed
- Univalent states
  - Outcome is fixed
  - May not be “known” yet

# Summary

---

- Wait-free computation is a tree
- Bivalent system states
  - Outcome not fixed
- Univalent states
  - Outcome is fixed
  - May not be “known” yet
- 1-Valent and 0-Valent states

# Claim

---

- Some initial state is bivalent

# Claim

---

- Some initial state is bivalent
- Outcome depends on
  - Chance
  - Whim of the scheduler

# Claim

---

- Some initial state is bivalent
- Outcome depends on
  - Chance
  - Whim of the scheduler
- Multicore gods procrastinate ...


# Claim

---

- Some initial state is bivalent
- Outcome depends on
  - Chance
  - Whim of the scheduler
- Multicore gods procrastinate ...
- Let's prove it ...

# What if inputs differ?

---







# Must Decide 0

In this solo execution by A



# Must Decide 1

---





In this solo execution by B


# Mixed Initial State Bivalent

- Solo execution by A must decide 0


- Solo execution by B must decide 1



# Critical States




# From a Critical State



If **A goes first**, protocol **decides 0**

If **B goes first**, protocol **decides 1**

# Reaching Critical State



# Critical States

---

- Starting from a bivalent initial state

# Critical States

---

- Starting from a bivalent initial state
- The protocol must reach a critical state



# Critical States

---

- Starting from a bivalent initial state
- The protocol must reach a critical state
  - Otherwise we could stay bivalent forever
  - And the protocol is not wait-free

# Model Dependency

---

- So far, memory-independent!
- True for
  - Registers
  - Message-passing
  - Carrier pigeons
  - Any kind of asynchronous computation

# Atomic Registers & Consensus

---

- Atomic registers have consensus number 1

# Closing the Deal

---

- Start from a critical state

# Closing the Deal

---

- Start from a critical state
- Each thread fixes outcome by
  - Reading or writing ...
  - Same or different registers

# Closing the Deal

---

- Start from a critical state
- Each thread fixes outcome by
  - Reading or writing ...
  - Same or different registers
- Leading to a 0 or 1 decision ...

# Closing the Deal

---

- Start from a critical state
- Each thread fixes outcome by
  - Reading or writing ...
  - Same or different registers
- Leading to a 0 or 1 decision ...
- **And a contradiction.**

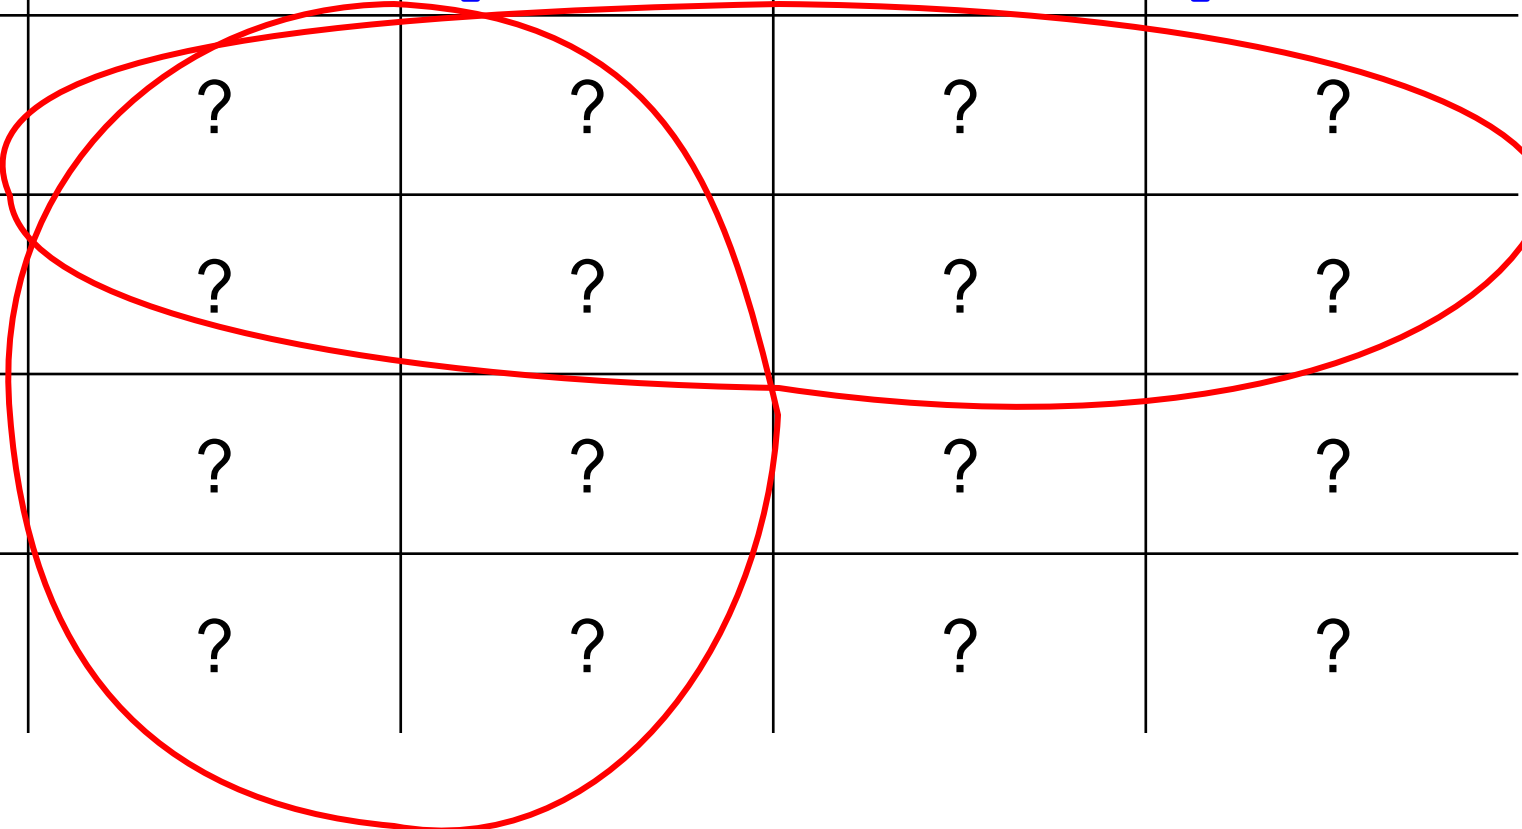
# Possible Interactions

	<code>x.read()</code>	<code>y.read()</code>	<code>x.write()</code>	<code>y.write()</code>
<code>x.read()</code>	<div><div>?</div><div>A reads x</div></div>	<div><div>?</div><div>A reads y</div></div>	?	?
<code>y.read()</code>	?	?	?	?
<code>x.write()</code>	?	?	?	?
<code>y.write()</code>	?	?	?	?




# Some Thread Reads

	<code>x.read()</code>	<code>y.read()</code>	<code>x.write()</code>	<code>y.write()</code>
<code>x.read()</code>	?	?	?	?
<code>y.read()</code>	?	?	?	?
<code>x.write()</code>	?	?	?	?
<code>y.write()</code>	?	?	?	?




# Some Thread Reads

---




# Some Thread Reads

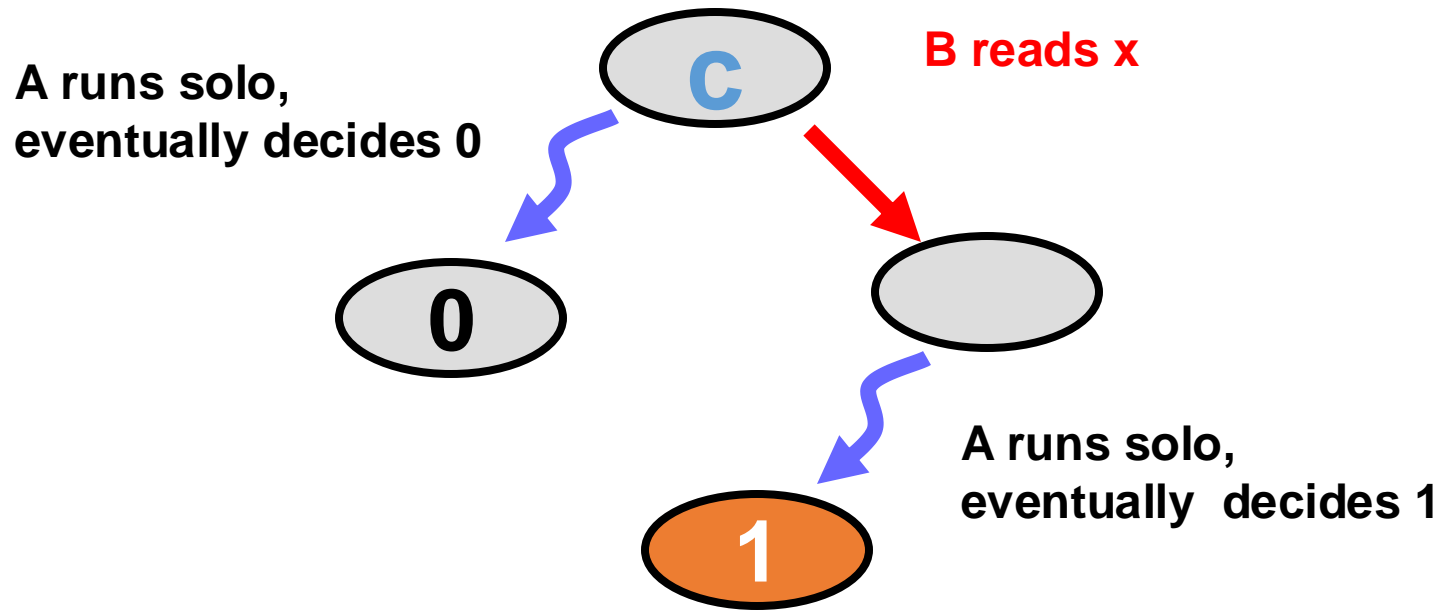
A runs solo,  
eventually decides 0



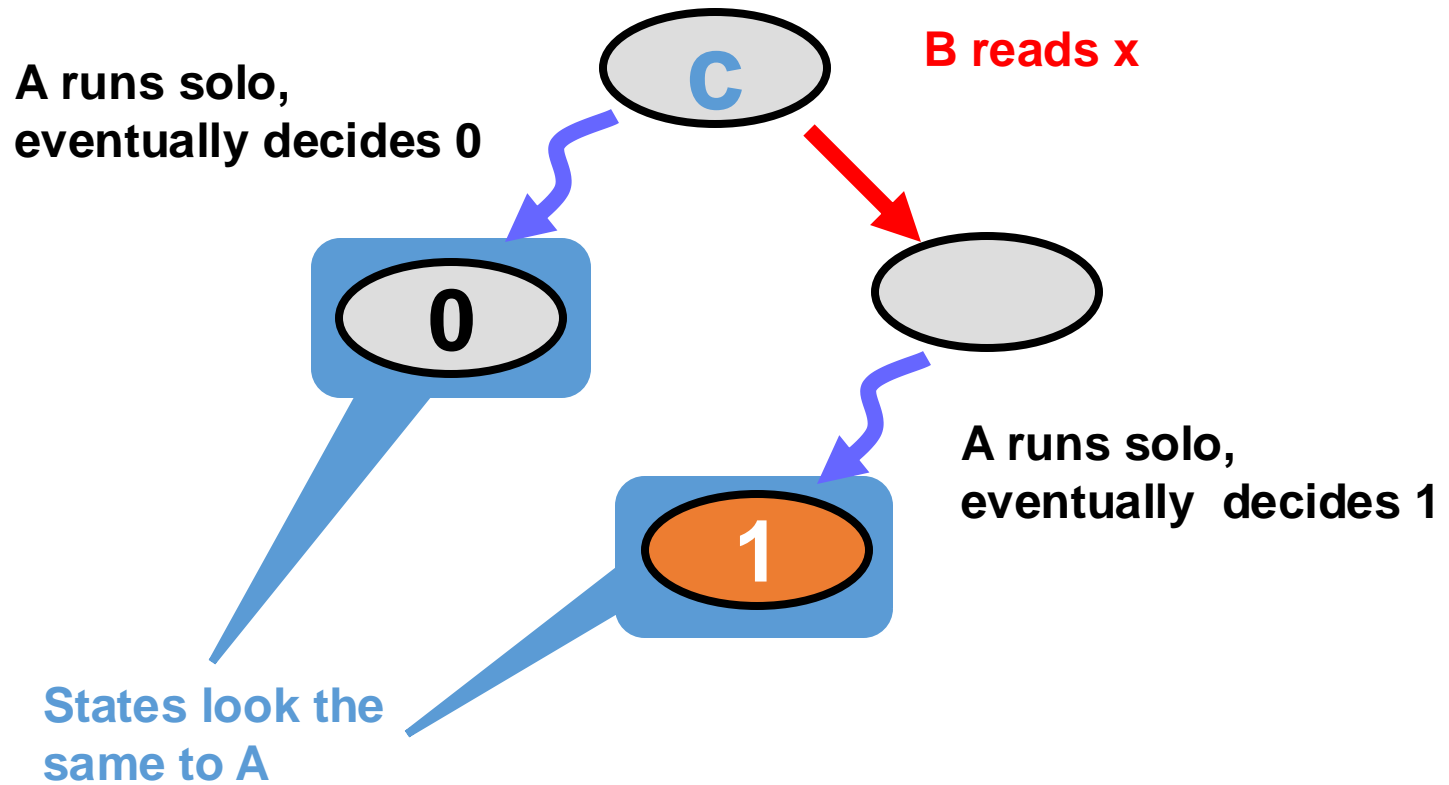
# Some Thread Reads




# Some Thread Reads



# Some Thread Reads



# Some Thread Reads



# Possible Interactions

	<code>x.read()</code>	<code>y.read()</code>	<code>x.write()</code>	<code>y.write()</code>
<code>x.read()</code>	no	no	no	no
<code>y.read()</code>	no	no	no	no
<code>x.write()</code>	no	no	?	?
<code>y.write()</code>	no	no	?	?




# Writing Distinct Registers


	<code>x.read()</code>	<code>y.read()</code>	<code>x.write()</code>	<code>y.write()</code>
<code>x.read()</code>	no	no	no	no
<code>y.read()</code>	no	no	no	no
<code>x.write()</code>	no	no	?	?
<code>y.write()</code>	no	no	?	?

# Writing Distinct Registers


---




# Writing Distinct Registers



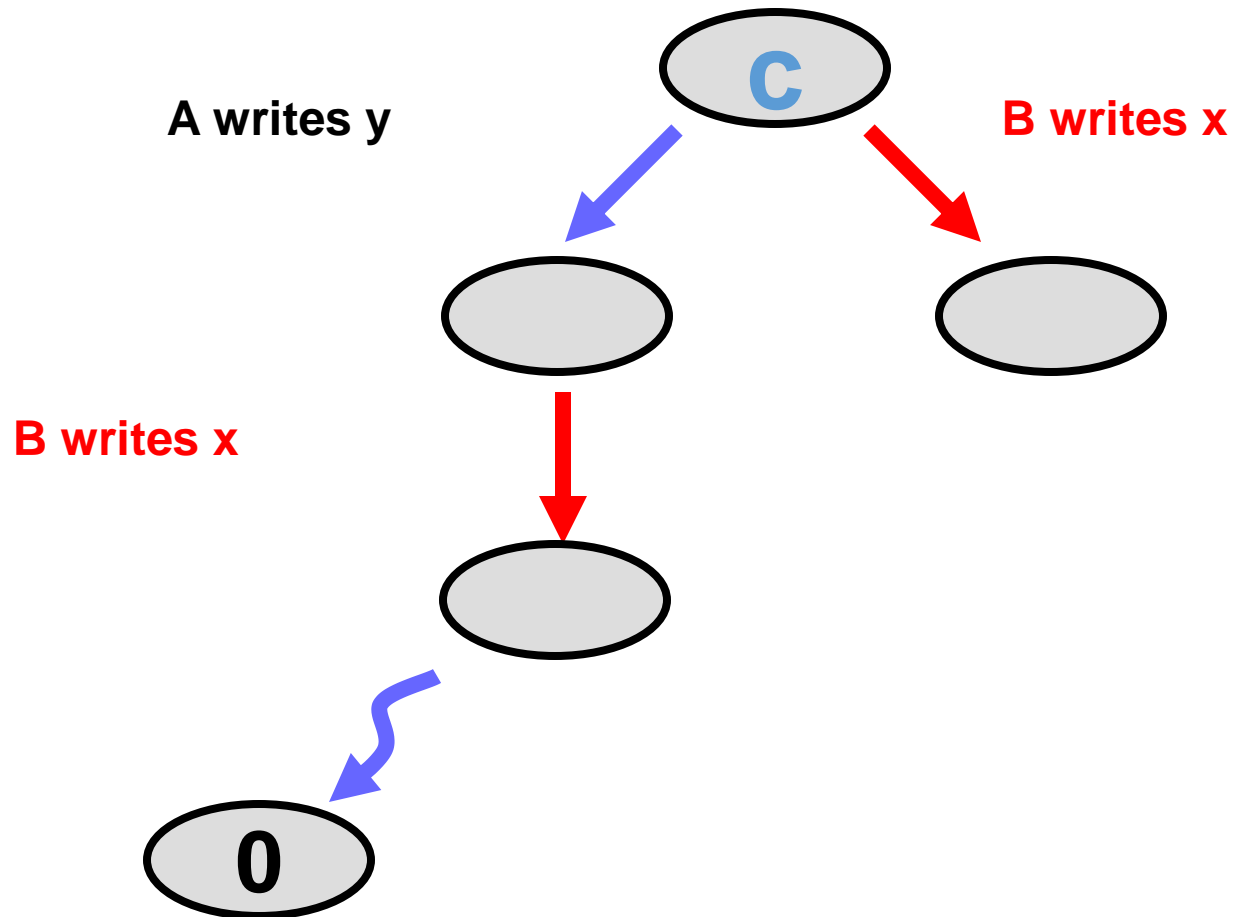
# Writing Distinct Registers




# Writing Distinct Registers



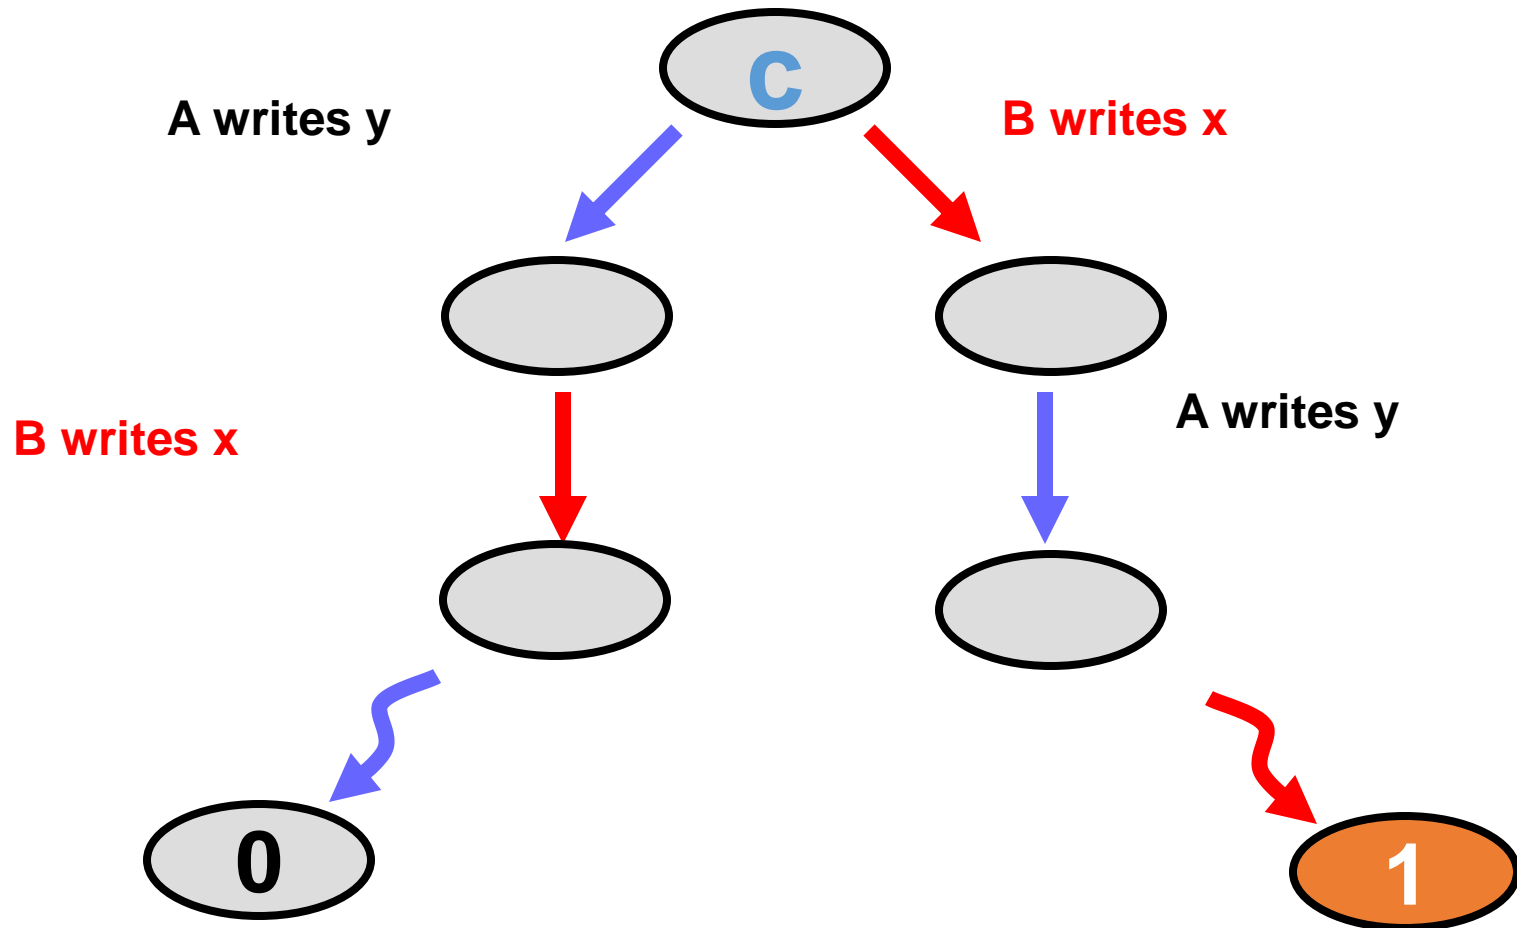
# Writing Distinct Registers



# Writing Distinct Registers

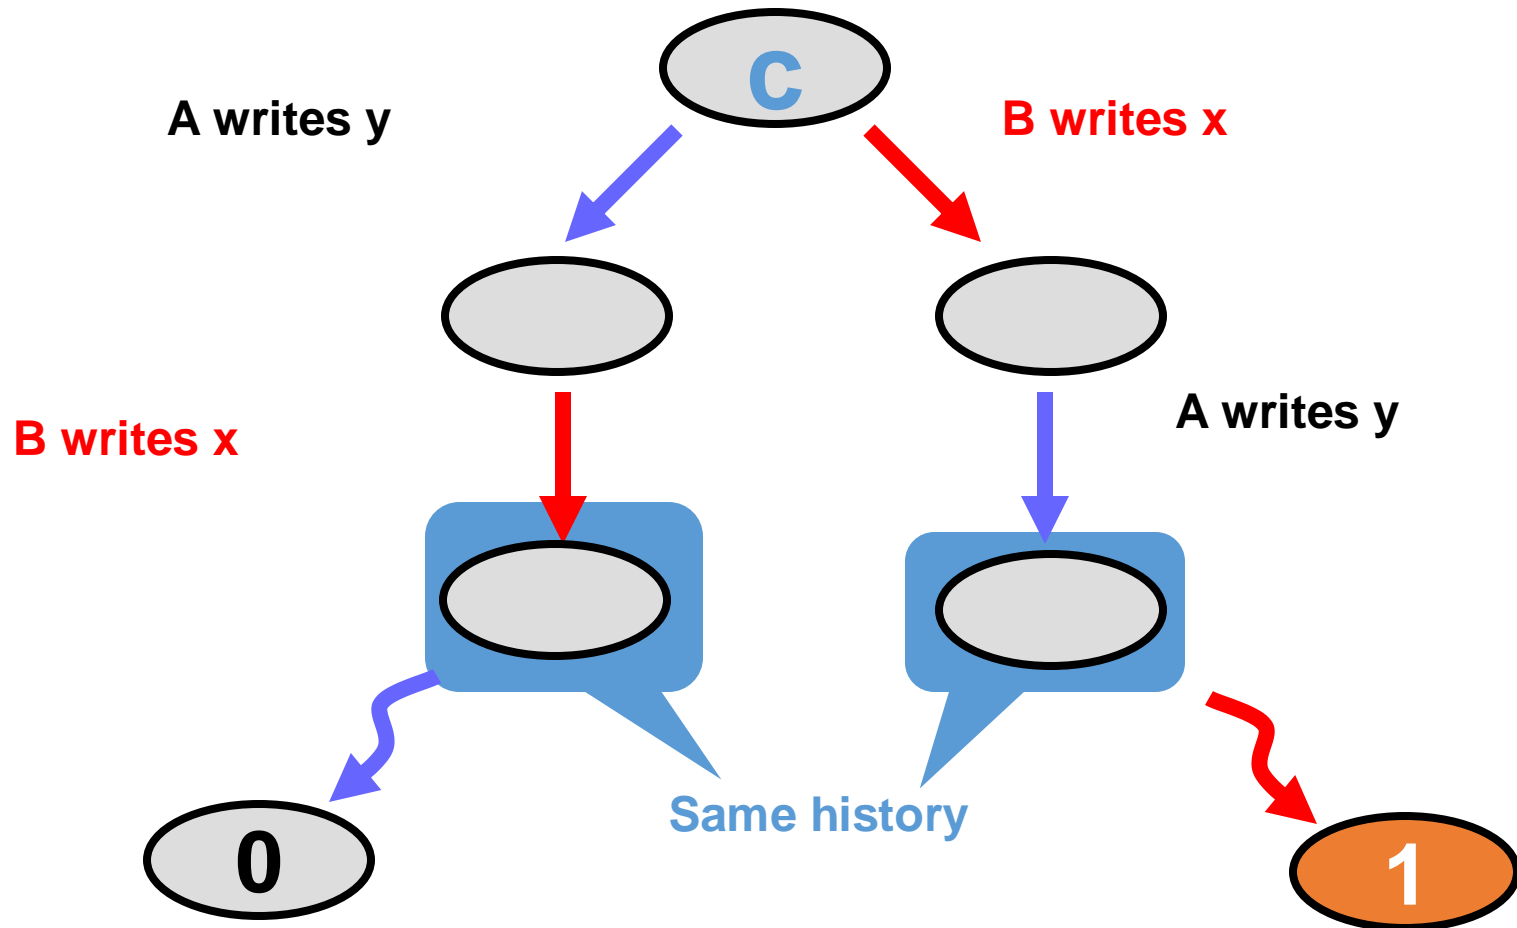


# Writing Distinct Registers






# Writing Distinct Registers



# Writing Distinct Registers




# Possible Interactions

	<code>x.read()</code>	<code>y.read()</code>	<code>x.write()</code>	<code>y.write()</code>
<code>x.read()</code>	no	no	no	no
<code>y.read()</code>	no	no	no	no
<code>x.write()</code>	no	no	?	no
<code>y.write()</code>	no	no	no	?

# Writing Same Registers


	<code>x.read()</code>	<code>y.read()</code>	<code>x.write()</code>	<code>y.write()</code>
<code>x.read()</code>	no	no	no	no
<code>y.read()</code>	no	no	no	no
<code>x.write()</code>	no	no	?	no
<code>y.write()</code>	no	no	no	?

# Writing Same Registers



# That's All, Folks!

	x.read()	y.read()	x.write()	y.write()
x.read()	no	no	no	no
y.read()	no	no	no	no
x.write()	no	no	no	no
y.write()	no	no	no	no




# Recap: Atomic Registers Can't Do Consensus

---

- If protocol exists
  - It has a bivalent initial state
  - Leading to a critical state
- What's up with the critical state?
  - Case analysis for each pair of methods
  - As we showed, all lead to a contradiction

# What Does Consensus have to do with Concurrent Objects?





# Consensus Object

---

```
public interface Consensus<T> {  
    T decide(T value);  
}
```

# Concurrent Consensus Object

---

- We consider only one time objects:
  - each thread calls method only once
- Linearizable to sequential consensus object:
  - Winner's call went first

# Java Jargon Watch

---

- Define Consensus protocol as an abstract class
- We implement some methods
- You do the rest ...

# Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>
    implements Consensus<T> {
    protected T[] proposed = new T[N];

    protected void propose(T value) {
        proposed[ThreadID.get()] = value;
    }

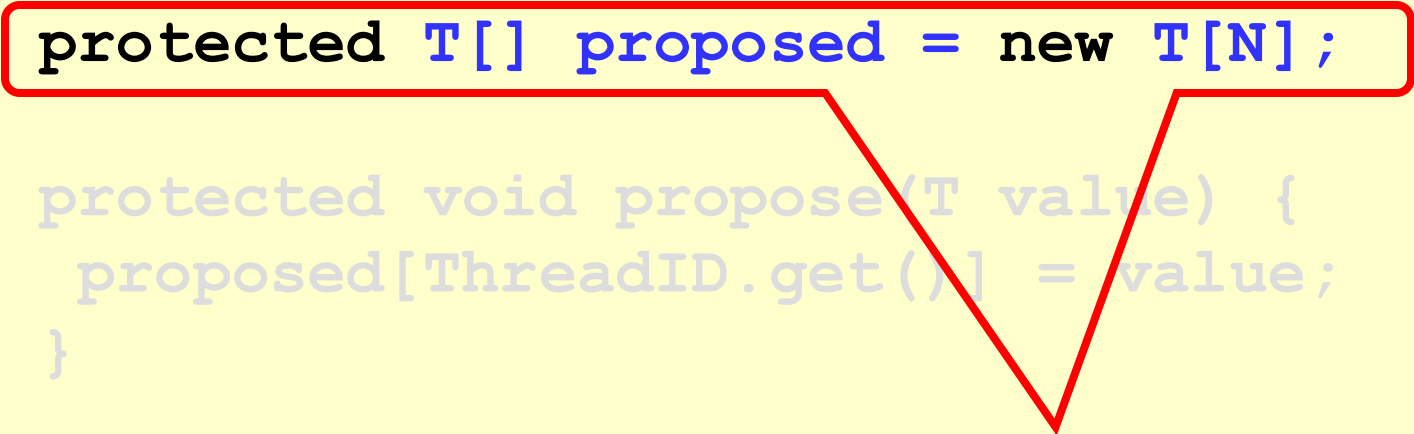
    abstract public T decide(T value);
}
```

# Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>
    implements Consensus<T> {
    protected T[] proposed = new T[N];

    protected void propose(T value) {
        proposed[ThreadID.get()] = value;
    }

    abstract public T decide()
}
```



Each thread's  
proposed value

# Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>  
    implements Consensus<T> {  
    protected T[] proposed = new T[N];
```

```
    protected void propose(T value) {  
        proposed[ThreadID.get()] = value;  
    }
```

```
    abstract public T decide(T value);  
}
```

**Propose a value**

# Generic Consensus Protocol

```
abstract class ConsensusProtocol<T>
    implements Consensus {
    protected T[] proposed = new T[N];
```

**Decide a value: abstract method means subclass  
does the real work**

```
    protected void propose(T value) {
        proposed[ThreadID.get()] = value;
    }
```


```
    abstract public T decide(T value);
```

---


# Can a FIFO Queue Implement Consensus?



# FIFO Consensus



**proposed array**






**Coveted red ball**

**Dreaded black ball**


**FIFO Queue  
with red and  
black balls**

# Protocol: Write Value to Array


---




# Protocol: Write Value to Array



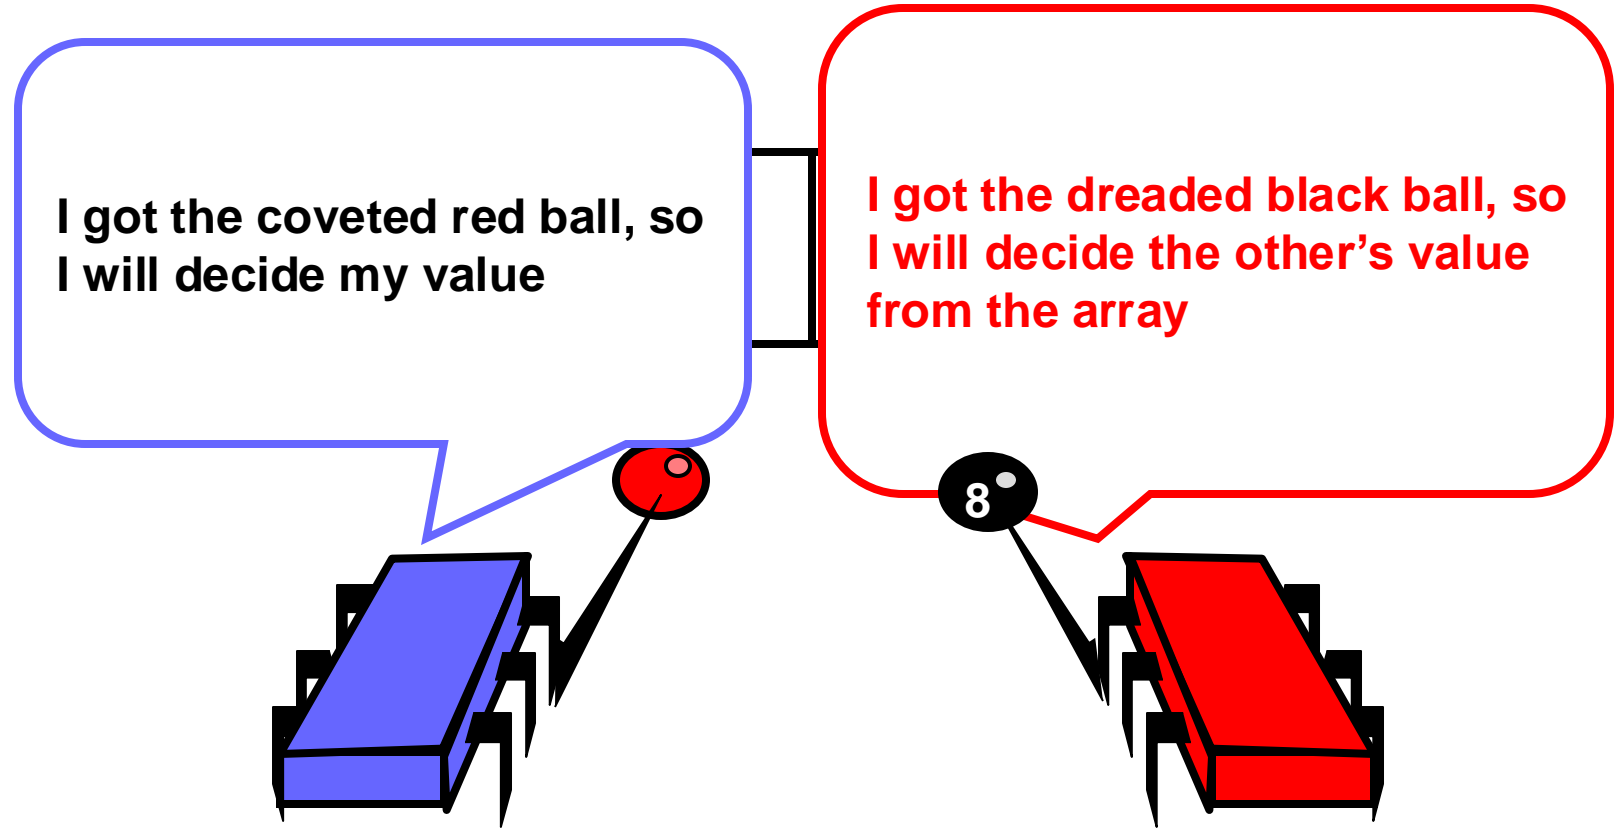
# Protocol: Write Value to Array



# Protocol: Take Next Item from Queue



# Protocol: Take Next Item from Queue




# Consensus Using FIFO Queue

```
public class QueueConsensus<T>
    extends ConsensusProtocol<T> {
    private Queue queue;
    public QueueConsensus () {
        queue = new Queue ();
        queue.enq(Ball.RED);
        queue.enq(Ball.BLACK);
    }
    ...
}
```

# Initialize Queue

```
public class QueueConsensus<T>
    extends ConsensusProtocol<T> {
    private Queue queue;
    public QueueConsensus() {
        queue = new Queue();
        queue.enq(Ball.RED);
        queue.enq(Ball.BLACK);
    }
    ...
}
```





# Who Won?

```
public class QueueConsensus<T>
    extends ConsensusProtocol<T> {
    private Queue queue;
    ...
    public T decide(T value) {
        propose(value);
        Ball ball = queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

# Who Won?

```
public class QueueConsensus<T>
    extends ConsensusProtocol<T> {
    private Queue queue;
    ...
    public T decide(T value) {
        propose(value);
        Ball ball = queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

**Race to dequeue first queue item**

# Who Won?

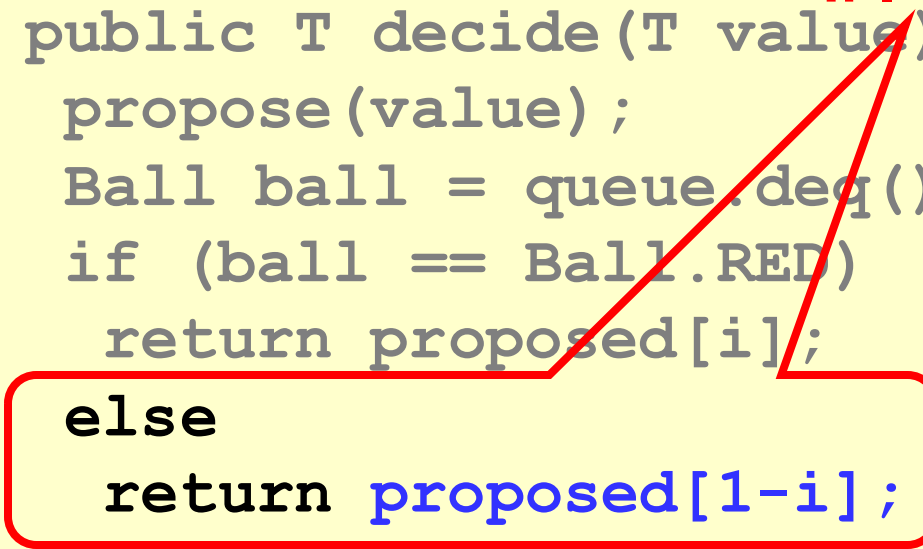
```
public class QueueConsensus<T>
    extends ConsensusProtocol<T> {
    private Queue queue;
    ...
    public T decide(T value) {
        propose(value);
        Ball ball = queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

**I win if I was first**

# Who Won?

```
public class QueueConsensus<T>
    extends ConsensusProtocol<T> {
    private Queue queue;
    ...
    public T decide(T value) {
        propose(value);
        Ball ball = queue.deq();
        if (ball == Ball.RED)
            return proposed[i];
        else
            return proposed[1-i];
    }
}
```

**Other thread wins  
if I was second**



# Why does this Work?

---

- If one thread gets the red ball
- Then the other gets the black ball
- Winner decides her own value
- Loser can find winner's value in array
  - Because threads write array
  - Before dequeuing from queue

# Theorem

---

- We can solve 2-thread consensus using only
  - A two-dequeueer queue, and
  - Some atomic registers

# Implications

- Given
  - A consensus protocol from queue and registers
- Assume there exists
  - A queue implementation from atomic registers
- Substitution yields:
  - A wait-free consensus protocol from atomic registers

**contradiction**

# Corollary

---

- It is impossible to implement
  - a two-dequeueer wait-free FIFO queue
  - from read/write memory.



# Consensus Numbers

---

- An object  $X$  has **consensus number  $n$** 
  - If it can be used to solve  $n$ -thread consensus
    - Take any number of instances of  $X$
    - together with atomic read/write registers
    - and implement  $n$ -thread consensus
  - But not  $(n+1)$ -thread consensus

# Consensus Numbers

---

- Theorem
  - Atomic read/write registers have consensus number 1
- Theorem
  - Multi-dequeueer FIFO queues have consensus number at least 2

# Consensus Numbers Measure Synchronization Power

---

- Theorem

- If you can implement  $X$  from  $Y$
- And  $X$  has consensus number  $c$
- Then  $Y$  has consensus number at least  $c$

# Synchronization Speed Limit

- Conversely
  - If  $X$  has consensus number  $c$
  - And  $Y$  has consensus number  $d$
  - Then there is no way to construct a wait-free implementation of  $X$  by  $Y$
- This theorem will be very useful
  - Unforeseen practical implications!

**Theoretical Caveat:  
Certain weird  
exceptions exist**

# Earlier Grand Challenge

---

- Snapshot means
  - Write any array element
  - Read multiple array elements atomically
- What about
  - Write multiple array elements atomically
  - Scan any array elements
- Call this problem **multiple assignment**

# Multiple Assignment Theorem

---

- Atomic registers cannot implement multiple assignment
- Weird or what?
  - Single write/multi read OK
  - Multi write/multi read impossible

# Proof Strategy

---

- If we can write to  $2/3$  array elements
  - We can solve 2-consensus
  - Impossible with atomic registers
- Therefore
  - Cannot implement multiple assignment with atomic registers


# Proof Strategy

---


- Take a 3-element array
  - A writes atomically to slots 0 and 1
  - B writes atomically to slots 1 and 2
  - Any thread can scan any set of locations




# Initially




# Thread **A** wins if



# Thread **A** wins if



# Thread **A** loses if



# Multi-Consensus Code

```
class MultiConsensus extends ...{
    Assign2 a = new Assign2(3, EMPTY);
    public T decide(T value) {
        a.assign(i, i, i+1, i);
        int other = a.read((i+2) % 3);
        if (other==EMPTY||other==a.read(1))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

# Multi-Consensus Code

```
class MultiConsensus extends ...{
    Assign2 a = new Assign2(3, EMPTY);
    public T decide(T value) {
        a.assign(i, i, i+1, i);
        int other = a.read((i+2) % 3);
        if (other==EMPTY||other==a.read(1))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

**Extends ConsensusProtocol**

**Decide sets  $j=i-1$  and proposes value**

# Multi-Consensus Code

```
class MultiConsensus extends ...{
    Assign2 a = new Assign2(3, EMPTY);
    public T decide(T value) {
        a.assign(i, i, i+1, i);
        int other = a.read((i+2) % 3);
        if (other==EMPTY||other==a.read(1))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

**Three slots  
initialized to  
EMPTY**

# Multi-Consensus Code

```
class MultiConsensus extends ...{
    Assign2 a = new Assign2(3, EMPTY);
    public T decide(T value) {
        a.assign(i, i, i+1, i);
        int other = a.read((i+2) % 3);
        if (other==EMPTY||other==a.read(1))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

**Assign ID 0 to entries 0,1  
(or ID 1 to entries 1,2)**



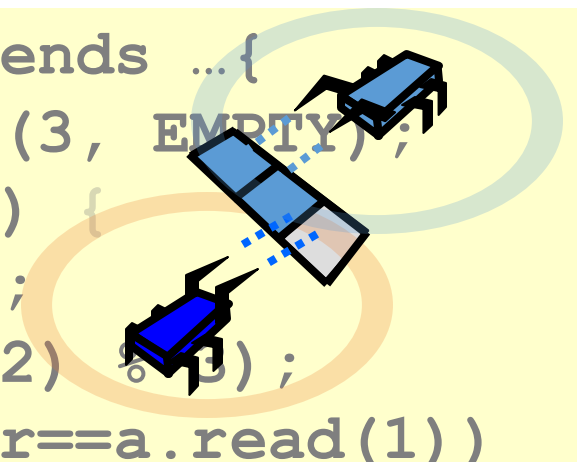
# Multi-Consensus Code

```
class MultiConsensus extends ...{
    Assign2 a = new Assign2(3, EMPTY);
    public T decide(T value) {
        a.assign(i, i, i+1, i);
        int other = a.read((i+2) % 3);
        if (other==EMPTY || other==a.read(1))
            return proposed[i];
        else
            return proposed[j];
    }
}
```

**Read the register my  
thread didn't assign**

# Multi-Consensus Code

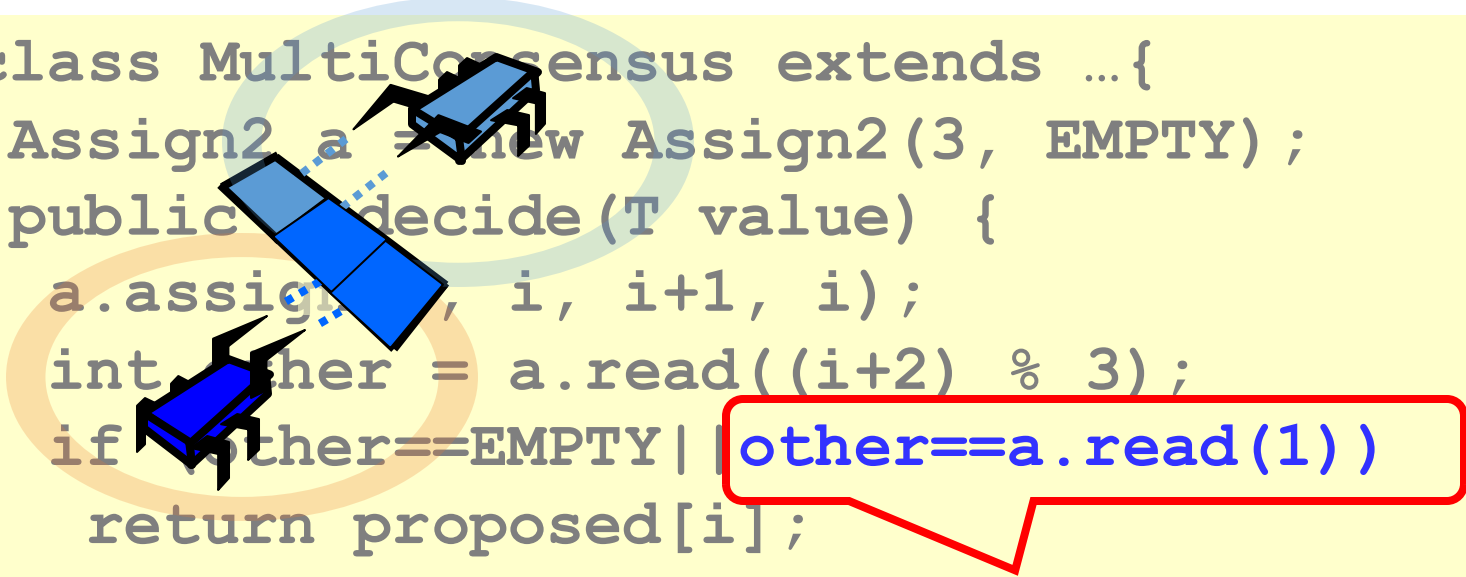
```
class MultiConsensus extends ...{
    Assign2 a = new Assign2(3, EMPTY);
    public T decide(T value) {
        a.assign(i, i, i+1, i);
        int other = a.read((i+2) % 3);
        if (other==EMPTY || other==a.read(1))
            return proposed[i];
        else
            return proposed[j];
    }
}
```



Other thread didn't move, so I win

# Multi-Consensus Code

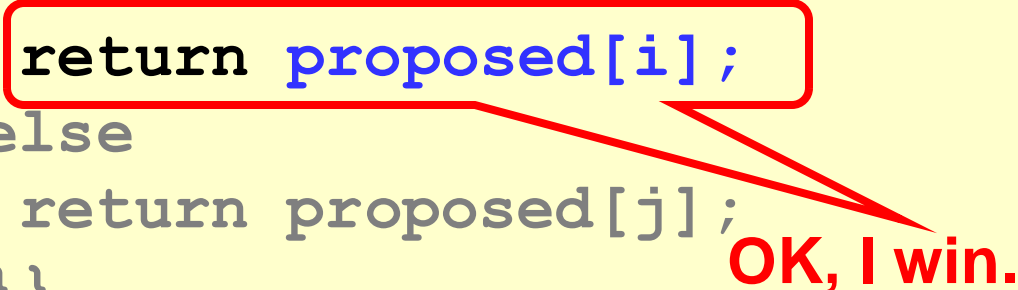
```
class MultiConsensus extends ...{
    Assign2 a = new Assign2(3, EMPTY);
    public T decide(T value) {
        a.assign(1, i, i+1, i);
        int other = a.read((i+2) % 3);
        if (other == EMPTY || other == a.read(1))
            return proposed[i];
        else
            return proposed[j];
    }
}
```



Other thread moved later so I win

# Multi-Consensus Code

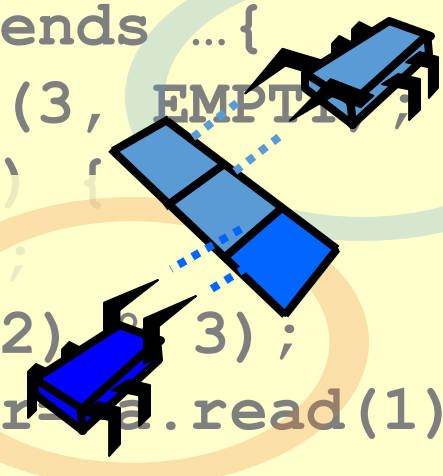
```
class MultiConsensus extends ...{
    Assign2 a = new Assign2(3, EMPTY);
    public T decide(T value) {
        a.assign(i, i, i+1, i);
        int other = a.read((i+2) % 3);
        if (other==EMPTY||other==a.read(1))
            return proposed[i];
        else
            return proposed[j];
    }
}
```



OK, I win.

# Multi-Consensus Code

```
class MultiConsensus extends ...{  
    Assign2 a = new Assign2(3, EMPTY);  
    public T decide(T value) {  
        a.assign(i, i, i+1, i);  
        int other = a.read((i+2), 3);  
        if (other==EMPTY || other==a.read(1))  
            return proposed[i];  
        else  
            return proposed[j];  
    }  
}
```



**Other thread moved  
first, so I lose**

# Summary

---

- If a thread can assign atomically to 2 out of 3 array locations
- Then we can solve 2-consensus
- Therefore
  - No wait-free multi-assignment
  - From read/write registers

# Read-Modify-Write Objects

---

- Method call
  - Returns object's prior value **x**
  - Replaces **x** with **mumble(x)**

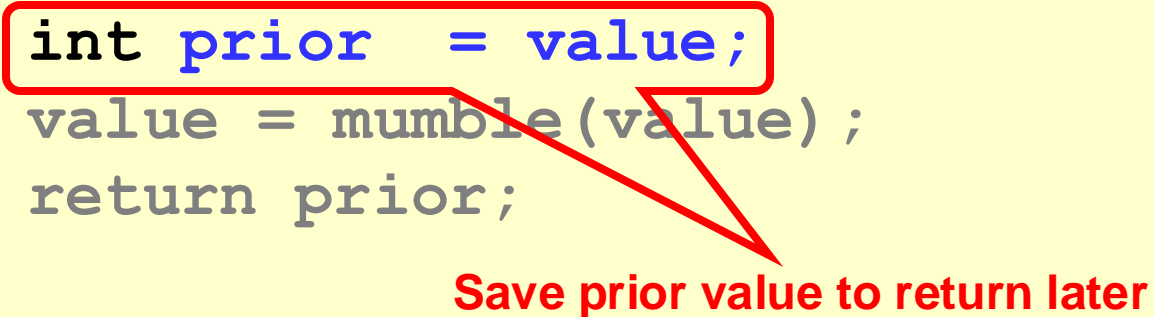
# Read-Modify-Write

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndMumble() {  
        int prior = value;  
        value = mumble(value);  
        return prior;  
    }  
}
```



# Read-Modify-Write

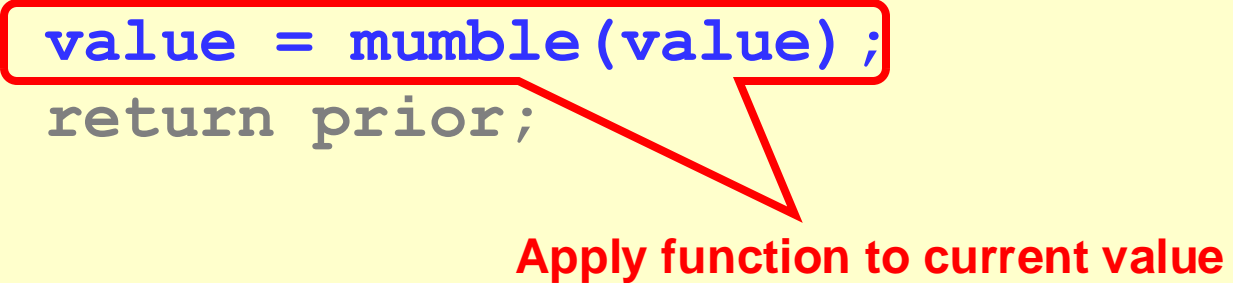
```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
    getAndMumble() {  
        int prior = value;  
        value = mumble(value);  
        return prior;  
    }  
}
```



Save prior value to return later

# Read-Modify-Write

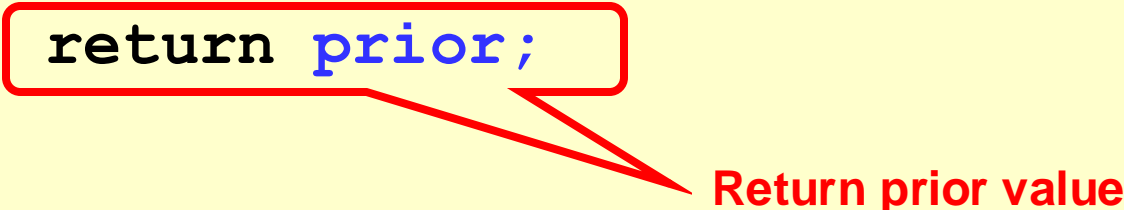
```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndMumble() {  
        int prior = value;  
        value = mumble(value);  
        return prior;  
    }  
}
```



**Apply function to current value**

# Read-Modify-Write

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndMumble() {  
        int prior = value;  
        value = mumble(value);  
        return prior;  
    }  
}
```



Return prior value

# RMW Everywhere!

---

- Most synchronization instructions
  - are RMW methods
- The rest
  - Can be trivially transformed into RMW methods

# Example: Read

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized read() {  
        int prior = value;  
        value = value;  
        return prior;  
    }  
}
```

# Example: Read

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized read() {  
        int prior = value;  
        value = value;  
        return prior;  
    }  
}
```

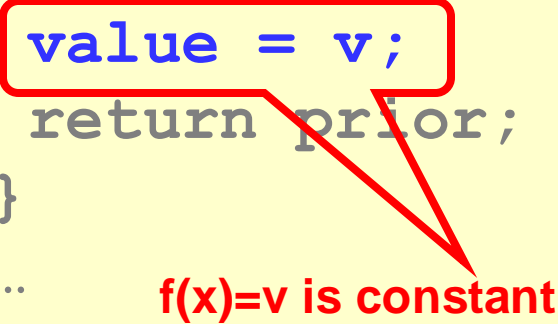
**apply  $f(x)=x$ , the identity function**

# Example: getAndSet (swap)

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndSet(int v) {  
        int prior = value;  
        value = v;  
        return prior;  
    }  
    ...  
}
```

# Example: getAndSet (swap)

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndSet(int v) {  
        int prior = value;  
        value = v;  
        return prior;  
    }  
    ...  
}
```



**f(x)=v is constant**

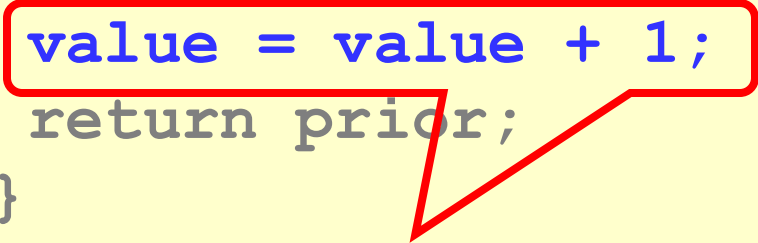


# getAndIncrement

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndIncrement() {  
        int prior = value;  
        value = value + 1;  
        return prior;  
    }  
    ...  
}
```

# getAndIncrement

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndIncrement() {  
        int prior = value;  
        value = value + 1;  
        return prior;  
    }  
    ...  
}
```



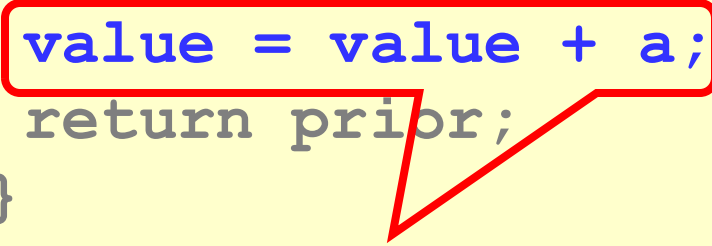
$f(x) = x+1$

# getAndAdd

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndAdd(int a) {  
        int prior = value;  
        value = value + a;  
        return prior;  
    }  
    ...  
}
```

# Example: getAndAdd

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndAdd(int a) {  
        int prior = value;  
        value = value + a;  
        return prior;  
    }  
    ...  
}
```



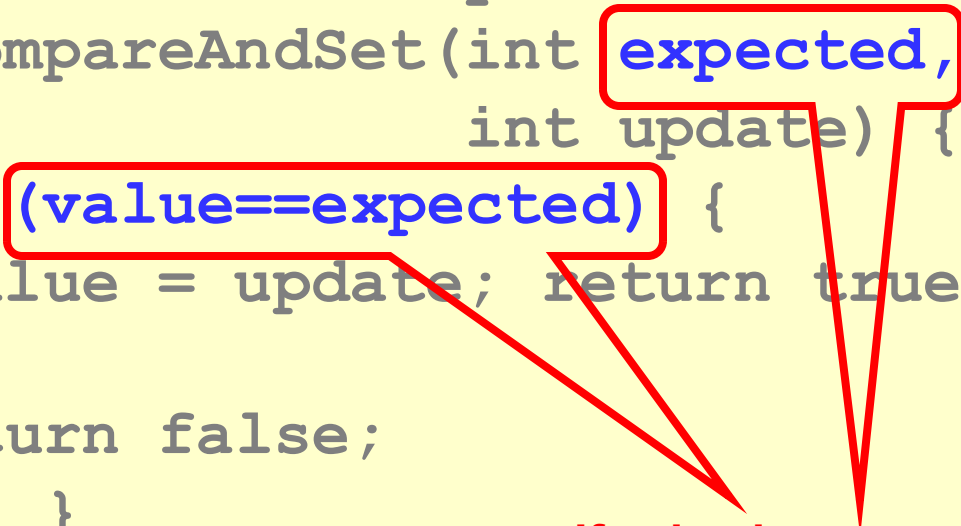
$f(x) = x+a$

# compareAndSet

```
public abstract class CASObject {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

# compareAndSet

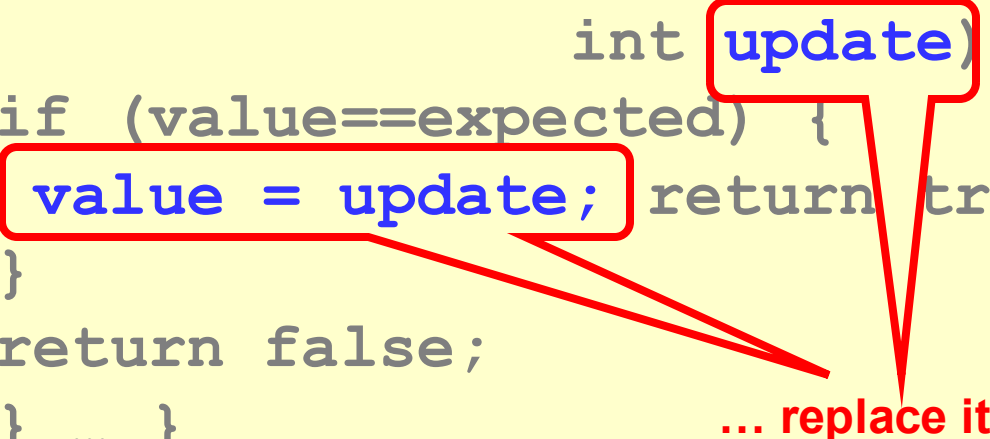
```
public abstract class CASObject {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

A diagram consisting of two red rectangular boxes. The first box is around the word 'expected' in the function signature of the 'compareAndSet' method. The second box is around the condition '(value==expected)' in the 'if' statement. Two red arrows originate from the bottom-right corner of the first box and the bottom-left corner of the second box, both pointing towards a red text note at the bottom right of the code block.

If value is as expected, ...

# compareAndSet

```
public abstract class CASObject {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  
    ... replace it
```



The diagram illustrates the replacement of the 'update' parameter with the current value of 'value' in the compareAndSet method. A red box highlights the 'update' parameter in the method signature. Another red box highlights the 'value = update;' assignment inside the if block. A red arrow points from the 'update' box to the 'value = update;' box, and another red arrow points from the 'value = update;' box to the text '... replace it'.

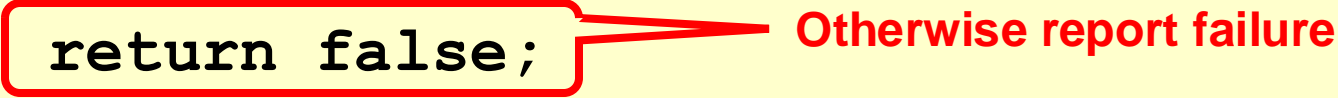
# compareAndSet

```
public abstract class CASObject {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

**Report success**

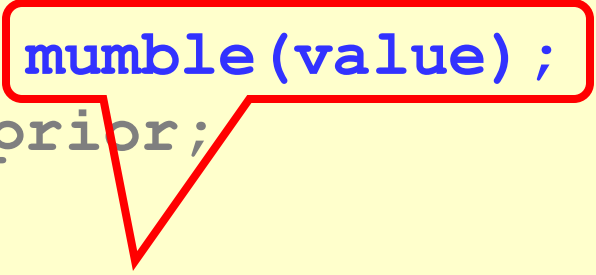


# compareAndSet

```
public abstract class CASObject {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;   
    } ... }  
}
```

# Read-Modify-Write

```
public abstract class RMWRegister {  
    private int value;  
  
    public int synchronized  
        getAndMumble() {  
        int prior = value;  
        value = mumble(value);  
        return prior;  
    }  
}
```



Let's characterize  $f(x)$ ...

# Definition

---

- A RMW method
  - With function `mumble(x)`
  - is non-trivial if there exists a value `v`
  - such that `v ≠ mumble(v)`

# Par Example

---

- **Identity(x) = x**
  - is trivial
- **getAndIncrement(x) = x+1**
  - is non-trivial

# Theorem

---

- Any non-trivial RMW object has consensus number at least 2
- No wait-free implementation of RMW registers from atomic registers
- Hardware RMW instructions not just a convenience

# Reminder

---

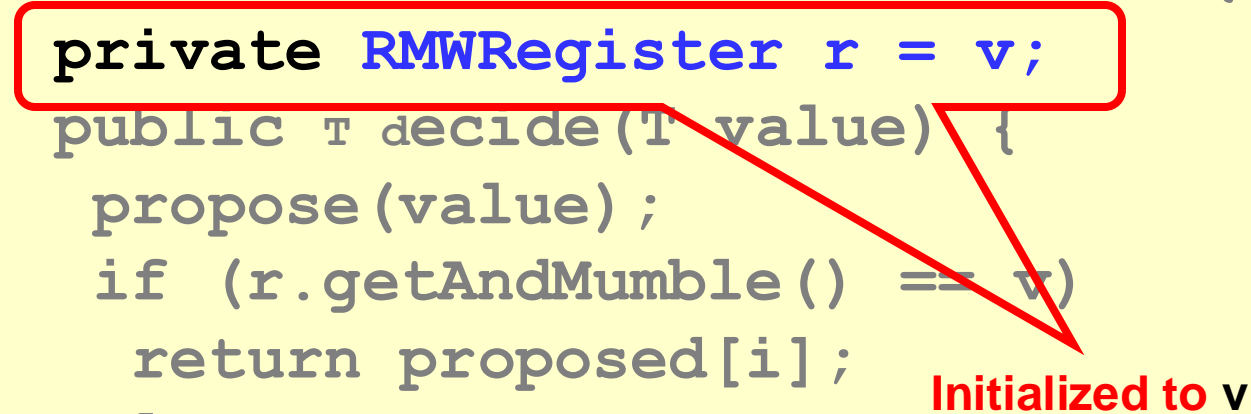
- Subclasses of **consensus** have
  - **propose(x)** method
    - which just stores x into **proposed[i]**
    - built-in method
  - **decide(object value)** method
    - which determines winning value
    - customized, class-specific method

# Proof

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = v;
    public T decide(T value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

# Proof

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = v;
    public T decide(T value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```



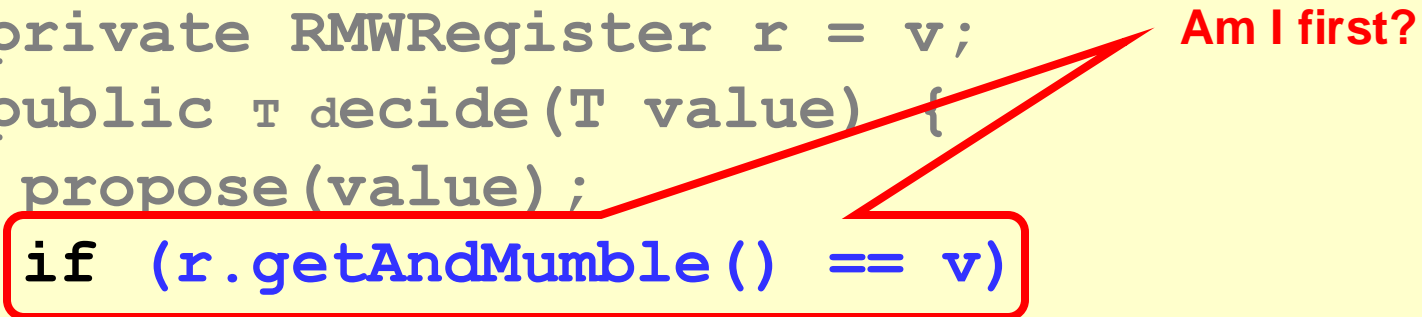
Initialized to v



# Proof

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = v;
    public T decide(T value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

Am I first?



# Proof

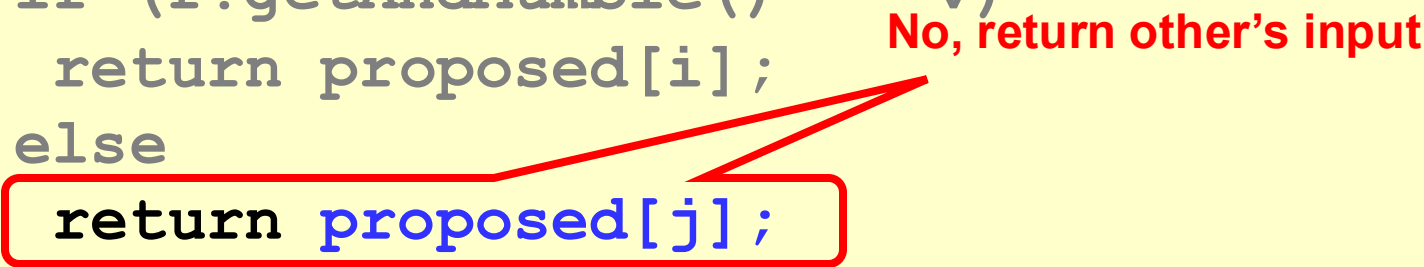
```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = v;
    public T decide(T value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

Yes, return my input

# Proof

```
public class RMWConsensus
    extends ConsensusProtocol {
    private RMWRegister r = v;
    public T decide(T value) {
        propose(value);
        if (r.getAndMumble() == v)
            return proposed[i];
        else
            return proposed[j];
    }
}
```

No, return other's input



# Proof

---

- We have displayed
  - A two-thread consensus protocol
  - Using any non-trivial RMW object

# Interfering RMW

- Let  $F$  be a set of functions such that for all  $f_i$  and  $f_j$ , either
  - Commute:  $f_i(f_j(v)) = f_j(f_i(v))$
  - Overwrite:  $f_i(f_j(v)) = f_i(v)$
- Claim: Any set of RMW objects that commutes or overwrites has consensus number exactly 2


# Examples

- “test-and-set”  $\text{getAndSet}(1)$   $f(v)=1$   
Overwrite  $f_i(f_j(v))=f_i(v)$
- “swap”  $\text{getAndSet}(x)$   $f(v,x)=x$   
Overwrite  $f_i(f_j(v))=f_i(v)$
- “fetch-and-inc”  $\text{getAndIncrement}()$   $f(v)=v+1$   
Commute  $f_i(f_j(v))=f_j(f_i(v))$


# Meanwhile Back at the Critical State

A about to apply  $f_A$

B about to apply  $f_B$

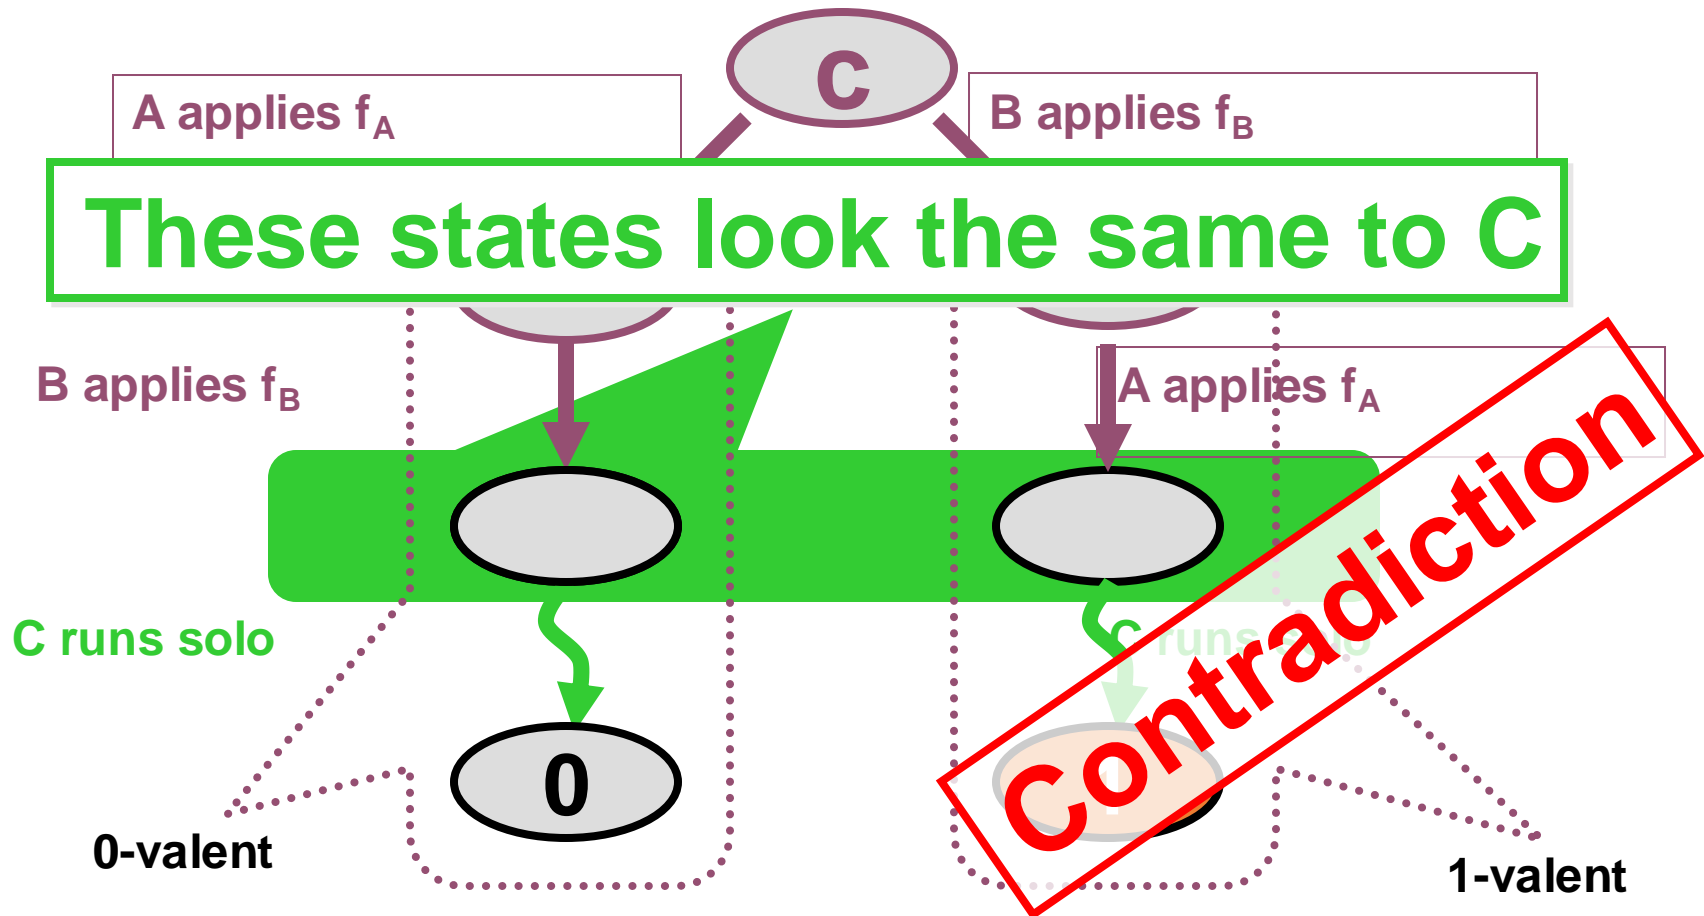


# Maybe the Functions Commute






# Maybe the Functions Commute




# Maybe the Functions Overwrite



# Maybe the Functions Overwrite

These states look the same to C



# Impact

---

- Many early machines provided only these “weak” RMW instructions
  - Test-and-set (IBM 360)
  - Fetch-and-add (NYU Ultracomputer)
  - Swap (Original SPARCs)
- We now understand their limitations
  - But why do we want consensus anyway?

# compareAndSet

```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = value;  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  
}
```

# compareAndSet

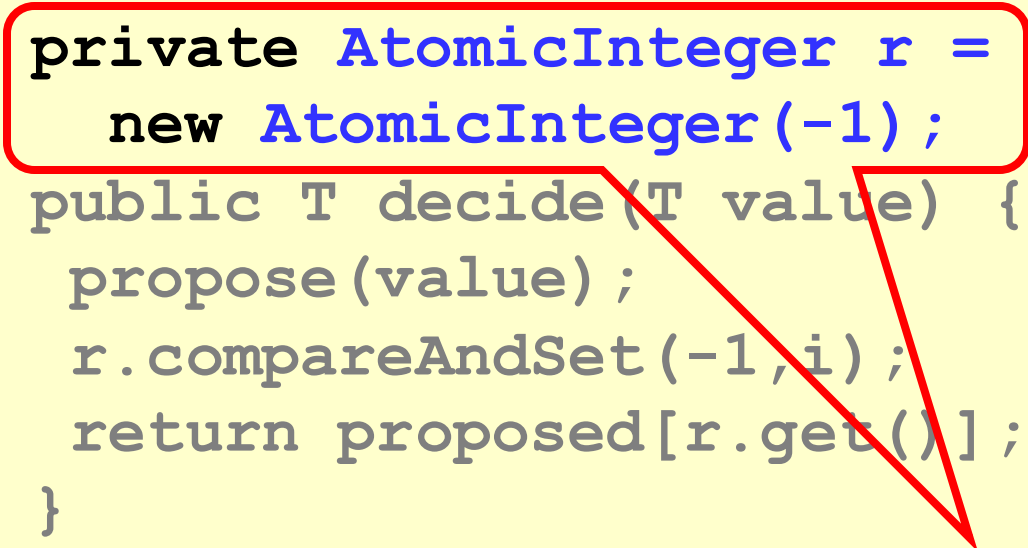
```
public abstract class RMWRegister {  
    private int value;  
    public boolean synchronized  
        compareAndSet(int expected,  
                       int update) {  
        int prior = value;  
        if (value==expected) {  
            value = update; return true;  
        }  
        return false;  
    } ... }  
    replace value if it's what we expected, ...
```

# compareAndSet Has $\infty$ Consensus Number

```
public class RMWConsensus
    extends ConsensusProtocol {
    private AtomicInteger r =
        new AtomicInteger(-1);
    public T decide(T value) {
        propose(value);
        r.compareAndSet(-1, i);
        return proposed[r.get()];
    }
}
```

# compareAndSet Has $\infty$ Consensus Number

```
public class RMWConsensus
    extends ConsensusProtocol {
    private AtomicInteger r =
        new AtomicInteger(-1);
    public T decide(T value) {
        propose(value);
        r.compareAndSet(-1, i);
        return proposed[r.get()];
    }
}
```



Initialized to -1



# compareAndSet Has $\infty$ Consensus Number

```
public class RMWConsensus
    extends ConsensusProtocol {
    private AtomicInteger r =
        new AtomicInteger(-1);
    public T decide(T value) {
        propose(value);
        r.compareAndSet(-1,i);
        return proposed[r.get()];
    }
}
```

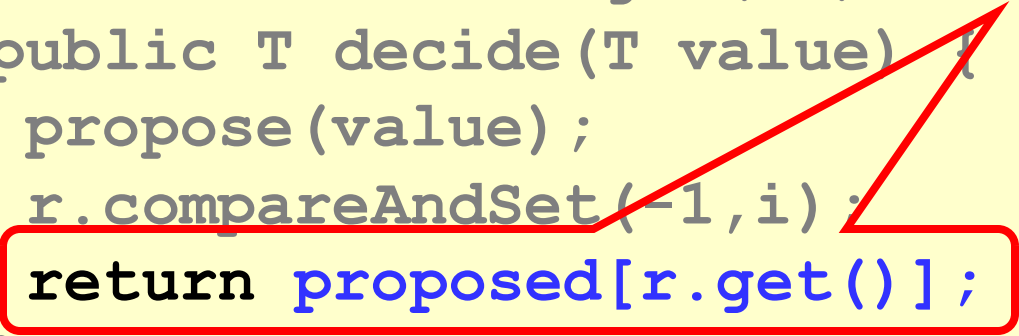
Try to swap in my id



# compareAndSet Has $\infty$ Consensus Number

```
public class RMWConsensus
    extends ConsensusProtocol {
    private AtomicInteger r =
        new AtomicInteger(-1);
    public T decide(T value) {
        propose(value);
        r.compareAndSet(-1, i);
        return proposed[r.get()];
    }
}
```

Decide winner's preference



# The Consensus Hierarchy

<b>1 Read/Write Registers, Snapshots...</b>
<b>2 getAndSet, getAndIncrement, ...</b>
<ul style="list-style-type: none"><li>▪</li><li>▪</li><li>▪</li></ul>
<b><math>\infty</math> compareAndSet,...</b>

# Lock-Freedom

---

- Lock-free:
  - in an infinite execution
  - infinitely often some method call finishes
- Pragmatic approach
- Implies no mutual exclusion



# Lock-Free vs. Wait-free

---

- Wait-Free: each method call takes a finite number of steps to finish
- Lock-free: infinitely often some method call finishes



# Lock-Freedom

- Any wait-free implementation is lock-free.
- Lock-free is the same as wait-free if the execution is finite.



# Lock-Free Implementations

---

- Lock-free consensus is as impossible as wait-free consensus
- ***All these results hold for lock-free algorithms also.***

# There is More: Universality

- Consensus is **universal**
- From  $n$ -thread consensus
  - Wait-free/Lock-free
  - Linearizable
  - $n$ -threaded
  - Implementation
  - Of any sequentially specified object

Stay tuned...



This work is licensed under a [Creative Commons Attribution-ShareAlike 2.5 License](https://creativecommons.org/licenses/by-sa/2.5/).

- **You are free:**
  - **to Share** — to copy, distribute and transmit the work
  - **to Remix** — to adapt the work
- **Under the following conditions:**
  - **Attribution.** You must attribute the work to “The Art of Multiprocessor Programming” (but not in any way that suggests that the authors endorse you or your use of the work).
  - **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under the same, similar or a compatible license.
- For any reuse or distribution, you must make clear to others the license terms of this work. The best way to do this is with a link to
  - <http://creativecommons.org/licenses/by-sa/3.0/>.
- Any of the above conditions can be waived if you get permission from the copyright holder.
- Nothing in this license impairs or restricts the author's moral rights.

# The END

---